AI EDITING BENCHMARK

How closely does each AI match your edits?

A side-by-side look at 3 editing tools measured against your own Lightroom adjustments.

TOOLS COMPARED

FotoLab, Imagen and Aftershoot

PHOTOS LEARNED FROM

5,409 photos

Image counts in this report have been anonymised: every training and test figure has been shifted by a small random amount and rounded away from its true value. The shape of every result is preserved, but the exact set sizes are not disclosed.

Basis of preparation

Every score in this report rests on one judgement call: how close is close enough? A prediction is counted as a hit only when it lands within a set distance of the value you actually chose - the **tolerance**. Each slider has its own, because a 200K shift in Temperature is barely perceptible while a 200-point shift in Contrast is significant.

These are the tolerances in force for this report. They drive both Ready to Deliver and Usable Sliders, and they are the dotted reference lines on the Average Accuracy charts.

SLIDER	TOLERANCE (PLUS OR MINUS)	UNITS
Temperature	200	K
Tint	4	pts
Exposure	0.25	stops
Contrast	5	pts
Highlights	10	pts
Shadows	5	pts
Whites	10	pts
Blacks	10	pts
Vibrance	4	pts
Saturation	4	pts

Tolerances are set under Metric Settings in the tool and saved with your preferences; this table records the values active when the report was generated.

What this report measures

Every tool here was given a batch of your photos to learn from, then asked to edit a separate batch on its own. Its edits on those photos were then compared with the edits you made yourself. No single number captures the whole picture, so the results are scored several different ways, each answering a different question:

Ready to Deliver

Of all the photos it edited, how many could you accept and deliver without touching a single slider?

Sliders on Target

Of every individual slider prediction, how many landed close enough that you would leave them alone?

Matched Your Style

How much better is the tool than blindly applying your average edit to every photo, the way a fixed preset would?

Average Accuracy

On average, how far off is each tool's slider, in the slider's own units? If you would have set Exposure to +0.50 and it predicted +0.30, that is a miss of 0.20 stops. This is the plain, traditional average error, with no normalising or scaling.

> **A NOTE ON HOW MUCH EACH TOOL LEARNED FROM**
>
> The scores below reflect a single batch - the number of your photos each tool learned from. The more of your work a tool sees, the better it should get at your style, so a result at one size is a snapshot rather than the whole story. Running the same tools on several batch sizes turns that snapshot into a curve, showing whether they keep improving with more of your photos or plateau early; it's worth testing more than one size before settling on a tool to live with.

Ready to Deliver

Higher is better.

WHAT THIS TELLS YOU

Of all the photos it edited, how many could you accept and deliver without touching a single slider?

An image only counts if every slider lands close enough; one bad prediction fails the whole image. A score of 30% means 300 of every 1,000 images can be delivered as they are.

> **TECHNICAL**
>
> This is the share of images on which every slider passes at once. It is the hardest metric to pass, because a single slider that marginally fails takes the whole image down with it.
>
> The effect compounds quickly. If every slider were individually accurate to:
>
> **85%** Only about 20% of images come out ready to deliver (0.85 to the tenth power).
>
> **90%** About 35%.
>
> **95%** About 60%.
>
> **99%** Still only about 90%.
>
> That steep curve is why this is almost always the lowest number on the board, it takes near-perfect per-slider accuracy to send most images out with no touch-ups at all.

Overall comparison

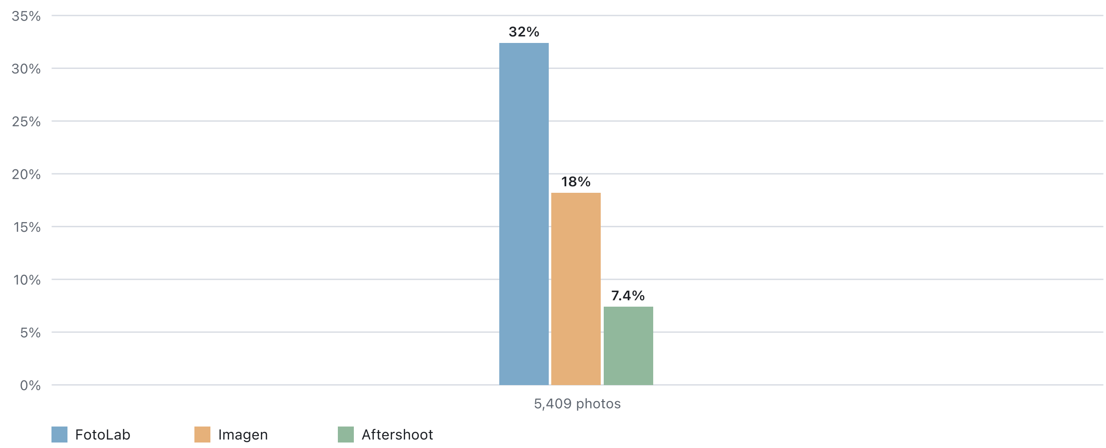

Sliders on Target

Higher is better.

WHAT THIS TELLS YOU

Of every individual slider prediction, how many landed close enough that you would leave them alone?

This counts sliders, not whole images. A score of 85% means 15 of every 100 predictions still need fixing.

> **TECHNICAL**
>
> This is the per-slider pass rate: the share of individual slider predictions that fall within tolerance. It counts individual sliders rather than whole images, so one photo with five bad sliders adds five failures here. That is the difference from Ready to Deliver, which marks that same photo as a single failure once any slider misses.

Overall comparison

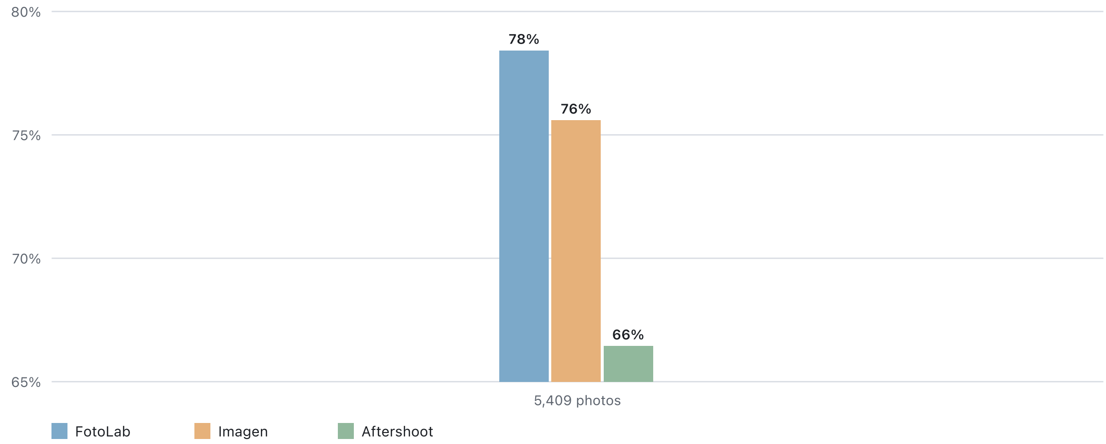

Score for each slider at 5,409 photos (the most photos every tool learned from)

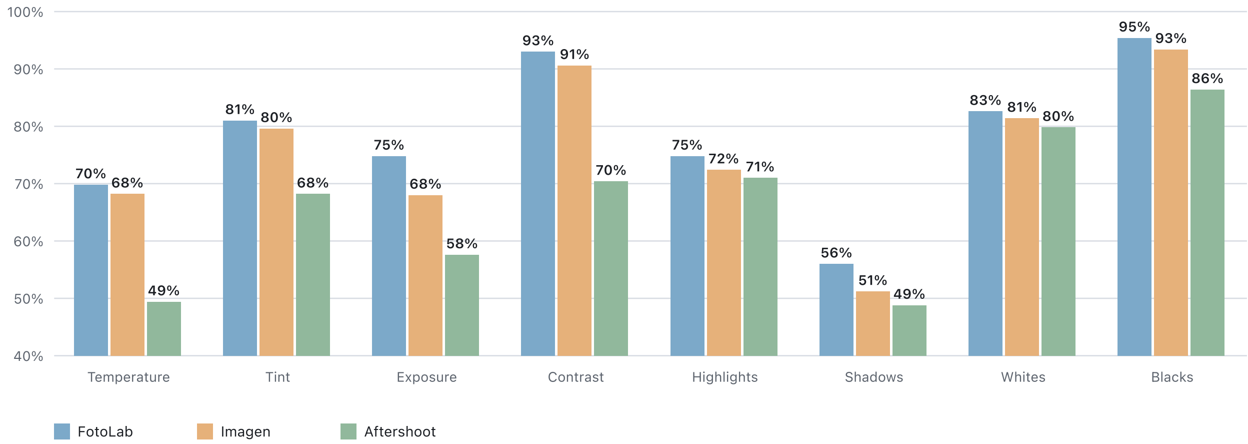

Matched Your Style

Higher is better.

WHAT THIS TELLS YOU

How much better is the tool than blindly applying your average edit to every photo, the way a fixed preset would?

100%	Reproduces your edits perfectly.
0%	No better than your flat average.
Negative	Worse than the average, which points to a fixed house style rather than yours.

Put simply, a high score means it has learned your style; a negative score means it is imposing its own.

> **TECHNICAL**
>
> This is the standard coefficient of determination (R^2). For each slider it weighs the tool's error against how much your own edits naturally vary, then averages those scores across all sliders. The yardstick it has to beat is your own average edit: the best possible one-size-fits-all version of you.
>
> Measuring each slider against its own variation makes them directly comparable. Temperature in Kelvin and Exposure in stops land on the same chart, and genuine differences between providers stay visible instead of being squashed toward 100% the way a simple range-based score would be.

Overall comparison

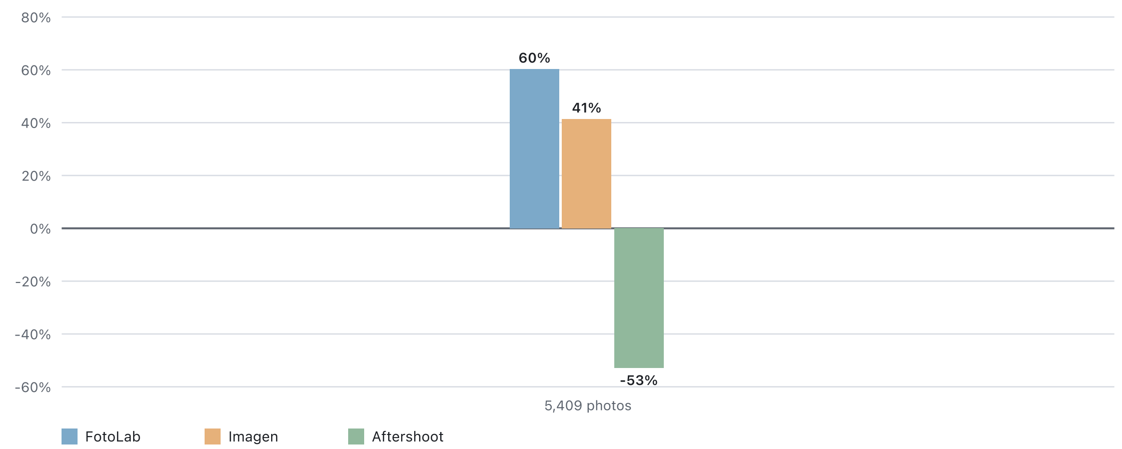

Score for each slider at 5,409 photos (the most photos every tool learned from)

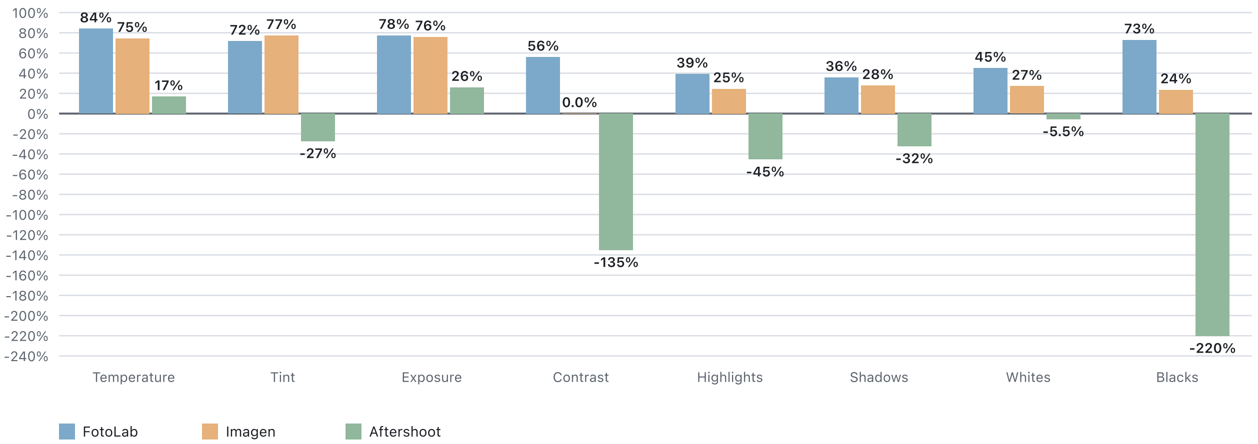

Average Accuracy

Taller is better (a smaller miss). The dotted line is the tolerance you set - a bar that rises above it is within tolerance.

WHAT THIS TELLS YOU

On average, how far off is each tool's slider, in the slider's own units? If you would have set Exposure to +0.50 and it predicted +0.30, that is a miss of 0.20 stops. This is the plain, traditional average error, with no normalising or scaling.

Each slider is anchored to the tolerance you set, drawn as a dotted line: a bar above the line is within tolerance, one below it misses. Taller is more accurate. Heights are scaled per slider, because a 200K Temperature miss and a 0.2-stop Exposure miss cannot share an axis; the number on each bar is the real average miss in native units.

> **TECHNICAL**
>
> This is the mean absolute error (MAE): the average of |predicted - your value| across all matched photos, computed per slider and kept in native units.
>
> Unlike R^2, nothing is divided by how much you moved the slider, so a slider you barely touched is not punished; the number is simply how far off the tool was, on average. For the same reason there is no single overall figure: averaging stops with Kelvin would be meaningless.

Score for each slider at 5,409 photos (the most photos every tool learned from)

Each slider has its own scale (Kelvin and points cannot share an axis). Taller is more accurate; a bar above the dotted tolerance line is within tolerance. Numbers are the real miss in native units.

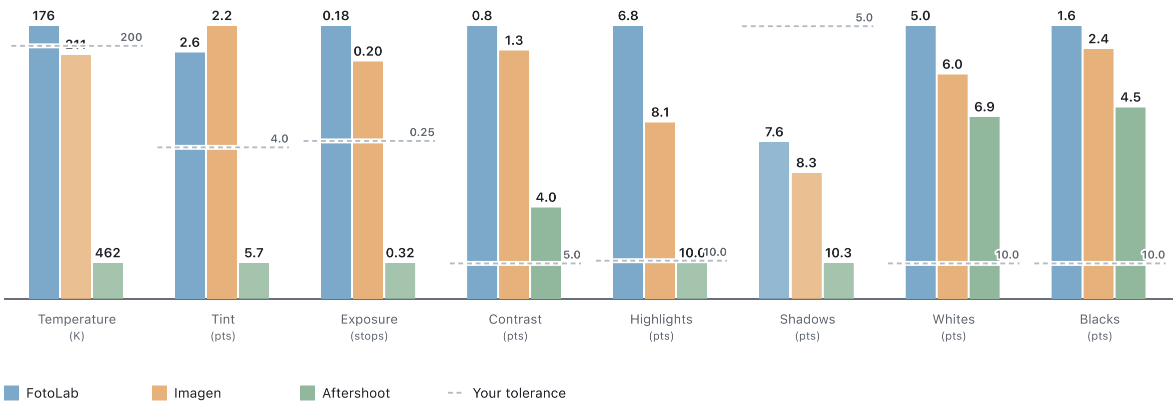

All the numbers

The table below lists every score behind the charts, broken down by tool and number of photos it learned from. Each metric runs from this report's earlier definitions.

Photos learned from	AI Tool	Matched	Ready to deliver whole images, no touch-ups	Sliders on target single sliders within tolerance	Matched your style vs blindly repeating your average
5,409 photos	FotoLab	595	32.4%	78.4%	60.3%
	Imagen	595	18.2%	75.6%	41.4%
	Aftershoot	599	7.4%	66.5%	−52.9%