

AI EDITING BENCHMARK

How closely does each AI match your edits?

A side-by-side look at 4 editing tools measured against your own Lightroom adjustments.

TOOLS COMPARED

FotoLab, Imagen, Aftershoot and Neaurapix

PHOTOS LEARNED FROM

1,189 photos, 4,168 photos, 10,125 photos

Image counts in this report have been anonymised: every training and test figure has been shifted by a small random amount and rounded away from its true value. The shape of every result is preserved, but the exact set sizes are not disclosed.

Basis of preparation

Every score in this report rests on one judgement call: how close is close enough? A prediction is counted as a hit only when it lands within a set distance of the value you actually chose - the **tolerance**. Each slider has its own, because a 200K shift in Temperature is barely perceptible while a 200-point shift in Contrast is significant.

These are the tolerances in force for this report. They drive both Ready to Deliver and Usable Sliders, and they are the dotted reference lines on the Average Accuracy charts.

SLIDER	TOLERANCE (PLUS OR MINUS)	UNITS
Temperature	200	K
Tint	4	pts
Exposure	0.25	stops
Contrast	5	pts
Highlights	10	pts
Shadows	5	pts
Whites	10	pts
Blacks	10	pts
Vibrance	4	pts
Saturation	4	pts

Tolerances are set under Metric Settings in the tool and saved with your preferences; this table records the values active when the report was generated.

What this report measures

Every tool here was given a batch of your photos to learn from, then asked to edit a separate batch on its own. Its edits on those photos were then compared with the edits you made yourself. No single number captures the whole picture, so the results are scored several different ways, each answering a different question:

Ready to Deliver

Of all the photos it edited, how many could you accept and deliver without touching a single slider?

Sliders on Target

Of every individual slider prediction, how many landed close enough that you would leave them alone?

Matched Your Style

How much better is the tool than blindly applying your average edit to every photo, the way a fixed preset would?

Average Accuracy

On average, how far off is each tool's slider, in the slider's own units? If you would have set Exposure to +0.50 and it predicted +0.30, that is a miss of 0.20 stops. This is the plain, traditional average error, with no normalising or scaling.

WHY THE AMOUNT EACH TOOL LEARNS FROM IS VARIED

A tool should match your style more closely the more of your photos it has learned from. Testing it once, with a single batch, shows only where it stands at that point. So each tool is tested several times - learning from, say, 2,000 of your photos, then 5,000, then 10,000 - and all edit the same photos at each step.

This is important to help understand real world usage. In normal use, images are continually added to a tool in the hope that it will become more accurate. This test reproduces that real-world usage and determines whether providers are actually learning, or whether they are failing to improve with increasing training sizes.

For every measure below, the first chart plots the score against the number of photos each tool learned from. The per-slider chart then compares the tools at the most photos they all learned from, so none is judged on a batch the others never edited.

Ready to Deliver

Higher is better.

WHAT THIS TELLS YOU

Of all the photos it edited, how many could you accept and deliver without touching a single slider?

An image only counts if every slider lands close enough; one bad prediction fails the whole image. A score of 30% means 300 of every 1,000 images can be delivered as they are.

TECHNICAL


This is the share of images on which every slider passes at once. It is the hardest metric to pass, because a single slider that marginally fails takes the whole image down with it.

The effect compounds quickly. If every slider were individually accurate to:


- 85%** Only about 20% of images come out ready to deliver (0.85 to the tenth power).
- 90%** About 35%.
- 95%** About 60%.
- 99%** Still only about 90%.

That steep curve is why this is almost always the lowest number on the board, it takes near-perfect per-slider accuracy to send most images out with no touch-ups at all.

How many photos you could deliver untouched, at each number of photos it learned from



How the score changes as it learns from more photos



Sliders on Target

Higher is better.


WHAT THIS TELLS YOU

Of every individual slider prediction, how many landed close enough that you would leave them alone?
This counts sliders, not whole images. A score of 85% means 15 of every 100 predictions still need fixing.


TECHNICAL

This is the per-slider pass rate: the share of individual slider predictions that fall within tolerance. It counts individual sliders rather than whole images, so one photo with five bad sliders adds five failures here. That is the difference from Ready to Deliver, which marks that same photo as a single failure once any slider misses.

How the score changes as it learns from more photos



Score for each slider at 10,125 photos (the most photos every tool learned from)



Matched Your Style

Higher is better.

WHAT THIS TELLS YOU

How much better is the tool than blindly applying your average edit to every photo, the way a fixed preset would?

- 100%** Reproduces your edits perfectly.
- 0%** No better than your flat average.
- Negative** Worse than the average, which points to a fixed house style rather than yours.


Put simply, a high score means it has learned your style; a negative score means it is imposing its own.

TECHNICAL


This is the standard coefficient of determination (R^2). For each slider it weighs the tool's error against how much your own edits naturally vary, then averages those scores across all sliders. The yardstick it has to beat is your own average edit: the best possible one-size-fits-all version of you.

Measuring each slider against its own variation makes them directly comparable. Temperature in Kelvin and Exposure in stops land on the same chart, and genuine differences between providers stay visible instead of being squashed toward 100% the way a simple range-based score would be.

How the score changes as it learns from more photos



Score for each slider at 10,125 photos (the most photos every tool learned from)



Average Accuracy

Taller is better (a smaller miss). The dotted line is the tolerance you set - a bar that rises above it is within tolerance.

WHAT THIS TELLS YOU

On average, how far off is each tool's slider, in the slider's own units? If you would have set Exposure to +0.50 and it predicted +0.30, that is a miss of 0.20 stops. This is the plain, traditional average error, with no normalising or scaling.

Each slider is anchored to the tolerance you set, drawn as a dotted line: a bar above the line is within tolerance, one below it misses. Taller is more accurate. Heights are scaled per slider, because a 200K Temperature miss and a 0.2-stop Exposure miss cannot share an axis; the number on each bar is the real average miss in native units.


TECHNICAL

This is the mean absolute error (MAE): the average of |predicted - your value| across all matched photos, computed per slider and kept in native units.

Unlike R^2 , nothing is divided by how much you moved the slider, so a slider you barely touched is not punished; the number is simply how far off the tool was, on average. For the same reason there is no single overall figure: averaging stops with Kelvin would be meaningless.

Score for each slider at 10,125 photos (the most photos every tool learned from)

Each slider has its own scale (Kelvin and points cannot share an axis). Taller is more accurate; a bar above the dotted tolerance line is within tolerance. Numbers are the real miss in native units.



All the numbers

The table below lists every score behind the charts, broken down by tool and number of photos it learned from. Each metric runs from this report's earlier definitions.

Photos learned from	AI Tool	Matched	Ready to deliver whole images, no touch-ups	Sliders on target single sliders within tolerance	Matched your style vs blindly repeating your average
10,125 photos	FotoLab	1,113	23.8%	75.3%	67.6%
	Imagen	1,113	8.1%	66.1%	49.0%
	Aftershoot	1,113	2.9%	59.8%	37.6%
	Neurapix	1,113	0.0%	40.3%	-7.0%
4,168 photos	FotoLab	1,113	11.5%	65.1%	50.6%
	Imagen	1,113	4.1%	60.8%	37.4%
	Aftershoot	1,113	0.7%	49.5%	17.8%
	Neurapix	1,113	0.0%	39.2%	-6.1%
1,189 photos	FotoLab	1,113	3.6%	54.7%	29.7%